

PPGEE2249 Aprendizado de Máquina
Assignment 3
Prof. Daniel Guerreiro e Silva

Thiago Tomás de Paula

November 26, 2025

► Question 1

Train and evaluate a Feedforward Neural Network (FNN) for a multiclass classification task. Split your data into a training set, a validation set (used to define the stopping criterion - e.g., number of epochs), and a test set for the final evaluation of the model. The output layer should use one Softmax neuron per class, and the training objective should be the minimization of cross-entropy loss. Describe your dataset, the network architecture (hidden units, layer dimensions, and number of layers), and discuss the overall performance (data splitting strategy, classification accuracy, and confusion matrix).

► Question 2

Train and evaluate a Decision Tree (DT) on the same classification problem chosen in Question 1. Present the resulting tree (depth, rule-based description) and compare its performance with the results obtained in Question 1.

Dataset description. The dataset chosen is Michael J. Fox Foundation’s Parkinson’s Progression Markers Initiative (PPMI) bank of raw sequencing reads from blood samples. Over 4000 participants have had their RNA transcribed across several visits, and the result of this analysis bundled with other categorical variables (such as age, gender, race, etc.) to form the PPMI dataset. As a dataset designed for the study of Parkinson Disease (PD) prediction and identification, each patient falls under one of three classes:

- “0”/“Case” if with PD
- “1”/“Control” if healthy
- “2”/“Other” if prodromal

Of the original 56 columns in the CSV of PPMI patient metadata (`meta_data11192021.csv`), 4 are diagnostic labels, and 32 were dropped from consideration due to one or more of the following:

- (i) redundancy (e.g., “Age (Bin)”, “Age at diagnosis”);
- (ii) could not contribute to the cause of PD (e.g., “Month” of visit, “Plate” of the sample);
- (iii) are patient identifiers (e.g., “HudAlphaSampleName”, “PATNO”);
- (iv) were sparsely present across samples (e.g., “Neutrophil Score”, UPDRS scores).

As such, each sample from the dataset consists of 20 features. From this subset of data, features with correlation greater than or equal to 90% (in absolute value) were dropped. E.g., if feature “A” showed correlation 0.93 with feature “B”, then either “B” or “A” is dropped, whichever is further at the right in the CSV file. Categorical features are first mapped to integers before calculating correlations.

After the removal of highly-correlated features, the samples are shuffled and finally separated into training, validation, and test sets in splits of 70%, 15% and 15%, respectively.

Table 1 profiles the sets obtained from the preprocessing of PPMI. Figure 1 shows a flowchart of the preprocessing done to data. Figure 2 shows the correlation matrix of the PPMI dataset. Since no correlation is bigger than 0.9 in absolute value, all 20 features were kept.

Table 1: Results from PPMI data preprocessing.

Class	Absolute and relative amounts within split					
	Train		Validation		Test	
0 (PD)	1352	58.83%	321	65.24%	300	60.85%
1 (Healthy)	843	36.68%	156	31.71%	172	34.89%
2 (Prodromal)	103	4.48%	15	3.05%	21	4.26%
Samples	2298	70.00%	492	14.99%	493	15.01%

Figure 1: Steps taken in this work to prepare data for machine learning algorithms.





Figure 2: Feature correlations found in a subset of PPMI.



Visit <https://www.ppmi-info.org/access-data-specimens/data> for more details on PPMI. Note that download of the metadata CSV requires applying for the data access.

FNN description. The FNN structure was obtained by minimizing validation error rate within a small search space of hyperparameter values. The hyperparameters considered and the associated search space boundaries are shown in table 2.

Table 2: Hyperparameters search space for FNN configuration

a	b
c	d

More precisely, the validation error corresponds to the number of wrong predictions made on the validation samples divided by the number of validation samples.

To find the point of minimum validation error within the hyperparameter space, the Optuna framework was used. This tool picks a subspace of the previously defined search space, samples hyperparameters from it, and evaluates the model using those values. The exploration of the subspace as a whole is called trial, and the model evaluations within a trial, steps. Trials are repeated arbitrarily, and their results are saved to guide following searches, with unpromising trials were stopped (pruned) early.

The best FNN (i.e., the one with smallest validation error) has the parameters presented in table 3. Figure 3 illustrates this network.

Table 3: FNN hyperparameters that minimize validation error.

a	b
c	d

Figure 3: FNN with smallest validation error

DT description. The DT's structure was obtained in the same fashion as the FNN's, that is, by minimizing validation error on a hyperparameter space. The hyperparameters considered and the associated search space boundaries are shown in table 4. The best FNN (i.e., the one with smallest validation error) has the parameters presented in table 5. Figure 4 illustrates this network.

Table 4: Hyperparameters search space for DT configuration

a	b
c	d


Table 5: DT hyperparameters that minimize validation error.

a	b
c	d

Performance comparison.

Figure 4: DT with smallest validation error

Figure 5: Confusion matrices of best FNN and DT architectures.

**► Question 3**

Perform an experiment testing different kernel functions and hyperparameter configurations to obtain the best SVM classifier for a binary classification task, using a dataset of your choice. Don't forget to discuss the results.
